
XData /sql

User Guide & Handbook

Automated SQL Query Grading Platform

A complete, screenshot-based guide for
Administrators, Instructors and Students

First start in 30 seconds: run `docker compose up --build`, open <http://localhost>,
and sign in with the seeded administrator account `admin1 / admin1`.

Contents

1	Introduction	2
2	Getting started	2
2.1	Starting the application	2
2.2	First-time initialisation	2
2.3	Signing in	2
3	Administrator guide	4
3.1	Creating users	4
3.2	Importing many users at once (CSV)	6
3.3	Managing existing users	6
3.4	Managing courses	8
4	Instructor guide	9
4.1	Defining schemas	10
4.2	Connecting a database server	11
4.3	Assignments and questions	13
4.4	Monitoring progress — the Gradebook	14
4.5	Other instructor tools	15
5	Student guide	18
5.1	The learning dashboard	18
5.2	Opening an assignment	19
5.3	Solving a question	20
5.4	Understanding the result	23
6	SQL Lab (Diagnostic Lab)	23
7	Dataset Playground	25
7.1	Mutation types	27
8	Tips & troubleshooting	28

1 Introduction

XData-Web is a web platform for learning SQL and grading SQL queries *automatically*. Instead of comparing query *text*, it checks whether a student's query is *semantically equivalent* to the instructor's reference solution, and — even when it is not — awards *partial credit* based on how close the query structure is (tables, columns, conditions, joins, grouping, ...).

This handbook walks through every part of the application, organised by the three user roles:

- **Administrator** — manages users and courses for the whole system.
- **Instructor** — creates courses, schemas, assignments and questions, and monitors student progress.
- **Student** — solves SQL questions and receives instant feedback.

Two additional tools are available to everyone who is signed in: the **SQL Lab** (a grading simulator) and the **Dataset Playground** (a test-data generator).

2 Getting started

2.1 Starting the application

The only prerequisite is **Docker**. From the project root, run:

```
docker compose up --build
```

This starts the database, the backend and the frontend. On the first start the database schema and a default administrator are created automatically, so there is no manual database setup. When the logs settle, open:

- **The app:** <http://localhost>
- **API docs (Swagger):** <http://localhost:8080/swagger-ui.html>

2.2 First-time initialisation

On a fresh installation exactly **one** account exists — the administrator. **Initialise the system by signing in with** `admin1 / admin1`.

Recommended first steps after the initial login:

1. Change the administrator password (key icon, top right).
2. Create a course (Admin Center → *Courses*).
3. Create at least one instructor and assign them to the course.
4. Let the instructor add a schema, an assignment and questions.
5. Create students and assign them to the course.

2.3 Signing in

Open the app and enter your user ID and password. New users receive their credentials from an administrator; the administrator starts with `admin1/admin1` (see §2.2).




Figure 1: The login screen. Sign in with your user ID and password.

The link “*Try SQL without signing in*” opens a public SQL sandbox — handy for trying the editor without an account.

A freshly created account is asked to **set its own password on first sign-in** before continuing — a security default you cannot skip.




Figure 2: First sign-in: new users must choose their own password before they can use the app.

3 Administrator guide

After signing in as an administrator you land in the **Admin Center**, which has two tabs: *Users* and *Courses*.

3.1 Creating users

The *Users* tab contains the **Create new user** form on the left and the list of existing users on the right.

The screenshot displays the 'Admin Center' interface for user management. At the top, there's a navigation bar with 'XData /sql', 'Dashboard', 'SQL Lab', and 'Dataset Playground'. The main header shows 'Admin Center' with 'System-wide management & monitoring.' and tabs for 'Users' and 'Courses'. Below this, the 'User management' section includes 'Template' and 'Import CSV' buttons. The 'Create new user' form contains fields for 'FULL NAME' (e.g. John Doe), 'LOGIN ID' (e.g. jdoe), 'EMAIL ADDRESS' (john@example.com), 'PASSWORD' (masked), and 'ROLE' (STUDENT). A 'COURSE ASSIGNMENT' section shows 'Datenbanken I' selected. A 'Create user' button is at the bottom right of the form. Below the form is a table of 'Assigned users' with 5 total users. The table has columns for STATUS, NAME, LOGIN ID, COURSE, ROLE, CREATED, and ACTIONS. The users listed are: schülerrr (STUDENT, 6/21/2026), Erika Lehrer (INSTRUCTOR, 6/20/2026), Max Schüler (STUDENT, 6/20/2026), Daniel Instructor (INSTRUCTOR, 6/17/2026), and Admin (ADMIN, 6/17/2026). Each row has a set of action icons (copy, delete, edit, share, trash).

ADMINISTRATION
Admin Center
System-wide management & monitoring.

USERS
User management

[Template](#) [Import CSV](#)

Create new user

FULL NAME: e.g. John Doe
LOGIN ID: e.g. jdoe
EMAIL ADDRESS: john@example.com
PASSWORD: *****
ROLE: STUDENT
COURSE ASSIGNMENT: ☒ Datenbanken I

[Create user](#)

Assigned users 5 TOTAL

Q Name, login ID or email... [Newest first](#)

STATUS	NAME	LOGIN ID	COURSE	ROLE	CREATED	ACTIONS
ACTIVE	S schülerrr	schülerrr	DB1-2026	STUDENT	6/21/2026	Copy Delete Edit Share Trash
ACTIVE	E Erika Lehrer	Lehrer1	DB1-2026	INSTRUCTOR	6/20/2026	Copy Delete Edit Share Trash
ACTIVE	M Max Schüler	schueler1	DB1-2026	STUDENT	6/20/2026	Copy Delete Edit Share Trash
ACTIVE	D Daniel Instructor	daniel	No course	INSTRUCTOR	6/17/2026	Copy Delete Edit Share Trash
ACTIVE	A Admin	admin1	No course	ADMIN	6/17/2026	Copy Delete Edit Share Trash

XData IIT Bombay · Automated SQL Grading © 2026 — built for learning, query by query.

Figure 3: Admin Center → Users: the user-creation form and the user list with per-row actions.

To create a single user:

1. Fill in **Full name**, **Login ID** (used to sign in) and **Email**.
2. Set an initial **Password**.
3. Choose a **Role**: STUDENT, INSTRUCTOR or ADMIN.
4. Tick one or more courses under **Course assignment**.
5. Click **Create user**. The new account appears in the list immediately.

The screenshot shows the 'Admin Center' interface for user management. The 'Create new user' form is the central focus, with fields for Full Name, Login ID, Email Address, Password, Role, and Course Assignment. The 'Assigned users' table below shows a list of users with columns for Status, Name, Login ID, Course, Role, Created, and Actions.

STATUS	NAME	LOGIN ID	COURSE	ROLE	CREATED	ACTIONS
ACTIVE	Anna Beispiel	anna1	Datenbanken I	INSTRUCTOR	2024-10-26	[Icons for actions]

Figure 4: The create-user form filled in for a new instructor, with a course ticked.

Roles in short: STUDENTS solve questions; INSTRUCTORS manage courses, schemas, assignments and grading; ADMINS manage all users and courses.

3.2 Importing many users at once (CSV)

For whole classes, use **Import CSV** (top right of the Users tab). Click **Template** first to download a correctly-formatted CSV, fill in one row per user, then upload it with **Import CSV**. All users are created in one step.

3.3 Managing existing users

Each row in the user list offers the following actions:

- **Disable / Enable** — block or restore sign-in without deleting the account.
- **Sign in as this user** — *impersonate* the user to see exactly what they see (great for support and for verifying an instructor's setup). A confirmation dialog appears first.
- **Edit** — change name, email, role or course assignment.
- **Reset password** — set a new password for the user.
- **Delete** — permanently remove the account.




Figure 5: The Edit-user dialog: change name, email, role and course assignment.




Figure 6: “Sign in as this user” asks for confirmation before impersonating — your admin session is replaced by the target user’s view.

3.4 Managing courses

The *Courses* tab lists all courses and lets you create new ones. A course bundles its instructors, students, schemas and assignments.




Figure 7: Admin Center → Courses: create and manage courses.

4 Instructor guide

Instructors work in the **Instructor Panel**. The left sidebar groups every tool; the *Advanced* section holds the setup tools (course materials, SQL schemas, databases, courses).




Figure 8: The Instructor Panel. The sidebar gives access to assignments, gradebook, students, statistics and the advanced setup tools.

4.1 Defining schemas

Before creating questions, define the database schema the questions run against. Open *Advanced* → *SQL Schemas*, select the course, and add the schema (table definitions). Schemas are reused by assignments, the SQL Lab and the Dataset Playground.




Figure 9: Schema management. Select a course, then create or edit the SQL schema used by its questions.

4.2 Connecting a database server

Besides the built-in structural and SMT checks, XData can grade by running both the student's query and the model query against a *real database server* and comparing the results — the most authoritative rung of the equivalence chain. To enable this for a course, register a **database connection** under *Advanced* → *Databases*.




Figure 10: Database connections for a course. Each card shows the connection name, course, JDBC URL and user.

Click **New Connection** and fill in four things plus the course:

- **Connection name** — a label, e.g. *Kurs-Server (PostgreSQL)*.
- **Course** — the course this server belongs to.
- **JDBC URL** — where the server is, in JDBC form (see below).
- **Username** and **Password** — the database credentials.




Figure 11: Creating a connection: name, course, JDBC URL, username and password.

Supported databases and JDBC URLs. XData ships with drivers for the following. Use the host *as the XData backend sees it* (for the bundled database that host is db; for an external server use its hostname or IP):

Database	JDBC URL format
PostgreSQL	jdbc:postgresql://host:5432/database
MySQL / MariaDB	jdbc:mysql://host:3306/database
SQLite	jdbc:sqlite:/path/to/file.db
H2	jdbc:h2:/path/to/db
Apache Derby	jdbc:derby:/path/to/db

Connections are verified on save. When you click **Save**, XData immediately opens a real connection to the server. If it cannot connect — wrong host, port, credentials or a firewall in the way — the connection is *rejected* with an error message, so every saved connection is known to work. For safety, XData also refuses to use its own *system* database as a course connection.

How it is used. A saved connection is bound to its course. When an assignment in that course is graded, XData can execute the queries against this server and compare the result sets (order-insensitive) — catching cases the cheaper checks leave undecided. You manage connections the same way later: edit to update the URL or credentials (re-verified on save), or delete one that is no longer needed.

4.3 Assignments and questions

Open *Assignments*, pick the course, and click **New** to create an assignment. For each assignment you set a **schema** and a **deadline**, and add **questions**. Every question has a problem statement, a

reference (model) solution, a difficulty, point value and optional hints. From an assignment you can also **Edit**, **Duplicate**, view **Statistics**, grant **Deadline extensions**, or **Delete** it. Set an assignment to *LIVE* to make it visible to students.




Figure 12: Creating an assignment: a two-step wizard — base configuration (name, deadline) then the question catalogue.

4.4 Monitoring progress — the Gradebook

The *Gradebook* shows every student's score per assignment, plus a total, and can be exported as **CSV**.

The screenshot shows the XData Instructor Panel interface. The top navigation bar includes 'XData /sql', 'Dashboard', 'SQL Lab', and 'Dataset Playground'. The user 'Erika Lehrer' is logged in as an instructor. The main section is titled 'INSTRUCTOR AREA' and 'Instructor Panel', with a subtitle 'Courses, assignments, and students in one place.'.

On the left is a sidebar menu with options: Assignments, Gradebook (highlighted), Regrade requests, Students, Announcements, Statistics, Course materials, SQL Schemas, Databases, and Courses.

The main content area is titled 'GRADING' and 'Gradebook'. It shows a table for the course 'Datenbanken I (DB1-2026)' with a 'CSV' export button. The table has three columns: 'STUDENT', 'WOCHE 1: SELECT BASICS', and 'TOTAL'.

STUDENT	WOCHE 1: SELECT BASICS	TOTAL
Max Schüler	10/10	100%
schülerrrr	0/10	0%

At the bottom, the footer reads 'XData IIT Bombay · Automated SQL Grading' and '© 2026 — built for learning, query by query.'

Figure 13: The Gradebook: per-student, per-assignment scores with a CSV export.

4.5 Other instructor tools

- **Regrade requests** — review and resolve student requests to re-check a submission.
- **Students** — see the roster of the selected course.
- **Announcements** — post messages to a course.
- **Statistics** — per-assignment analytics (open them via the chart icon on an assignment).
- **Databases** — register external database connections used for grading.
- **Course materials** — attach links and resources to a course.


XData /sql

[Dashboard](#)
[SQL Lab](#)
[Dataset Playground](#)




Anna Beispiel
INSTRUCTOR




INSTRUCTOR AREA
Instructor Panel
Courses, assignments, and students in one place.

[Assignments](#)
[Gradebook](#)
[Regrade requests](#)
[Students](#)
[Announcements](#)
[Statistics](#)

ADVANCED

COURSE PARTICIPANTS
User management

[Template](#)
[Add existing](#)
[Import CSV](#)


Create new student

FULL NAME

LOGIN ID

EMAIL ADDRESS

PASSWORD

COURSE ASSIGNMENT (MY COURSES)

☒ **Datenbanken I**

[Create student](#)

Assigned users
4 TOTAL

[Newest first](#)

STATUS	NAME	LOGIN ID	COURSE	ROLE	CREATED	ACTION
 ACTIVE	 A Anna Beispiel	anna1	DB1-2026	INSTRUCTOR	6/30/2026	  
 ACTIVE	 S schülerrr	schülerrr	DB1-2026	STUDENT	6/21/2026	  
 ACTIVE	 E Erika Lehrer	lehrer1	DB1-2026	INSTRUCTOR	6/20/2026	  
 ACTIVE	 M Max Schüler	schueler1	DB1-2026	STUDENT	6/20/2026	  

XData IIT Bombay · Automated SQL Grading

© 2026 — built for learning, query by query.

Figure 14: Students: the roster of the selected course.




Figure 15: Announcements: post a message (title, course, content) to a course; published announcements are listed below.




Figure 16: Regrade requests: review and resolve students' requests to re-check a submission.

5 Student guide

5.1 The learning dashboard

After signing in, a student sees their personal dashboard: level and XP progress, a daily streak, a **leaderboard**, earned **badges**, and the list of **assignments**.




Figure 17: The student dashboard: XP and level, leaderboard, badges and assigned modules.

5.2 Opening an assignment

Click an assignment under *My assignments* to load its questions. Each question shows its difficulty, tags and point value; solved questions are marked with a green check.




Figure 18: Selecting an assignment loads its question list on the right.

5.3 Solving a question

Open a question to reach the solving view. It contains:

- The **problem statement**, plus **Schema**, **Cheat sheet** and **Sample data** helpers.
- Step-by-step **hints** (optional — revealing them may cost points).
- The **SQL editor** with syntax highlighting and auto-completion.
- **Run** (try the query) and **Submit answer & check** (grade it).

The screenshot displays the XData-Web student interface. At the top, a dashboard shows the user's name 'Max Schüler', level 'Level 6', and progress '2544 / 3400 XP'. It also indicates '0 days streak' and 'Today 9/3'. The main content area is divided into several sections:

- LEADERBOARD:** Shows 'You' with 1 solved problem and 10 pts, and 'schülerrr' with 0 solved and 0 pts.
- BADGES:** Includes 'First Solution', 'Diligent', 'SQL Pro', 'Perfectionist', 'Climber', and 'On a Roll'.
- MY ASSIGNMENTS:** Shows 'Woche 1: SELECT Basics' with a 100% completion rate.
- QUESTIONS:** Displays 'Studierende über 22' (EASY, 10 pts) with a 'View history' link.
- PROBLEM STATEMENT:** 'Gib die Namen aller Studierenden aus, die älter als 22 Jahre sind.'
- HINTS:** 'Stuck? Reveal hints step by step.' with a 'Show hint (0/2)' button.
- SQL EDITOR:** A large text area for writing the SQL query.
- Buttons:** 'Run', 'Submit answer & check', 'Schema', 'Cheat sheet', and 'Sample data'.
- LAST RESULT:** Shows an 'OVERALL SCORE' of 3.0 / 3.0 (100%) and a 'MAIN QUERY' table with component feedback.

COMPONENT	YOUR SOLUTION	MODEL SOLUTION
Tables (FROM)	✓ STUDENTS (+1.0)	STUDENTS
Columns (SELECT)	✓ name (+1.0)	name
Conditions (WHERE)	✓ age > 22 (+1.0)	age > 22
Joins (JOIN)	-	-
Grouping (GROUP BY)	-	-
Sorting (ORDER BY)	-	-
Filter (HAVING)	-	-
Duplicates (DISTINCT)	✓ No (-0.0)	NO

At the bottom, there are links for 'Compare expected output' and 'Request regrade'.

Figure 19: The question view: problem statement, hints, SQL editor, and the structural result with per-component feedback.

Helpers. The **Schema** button reveals the tables, columns, primary keys and relationships of the question's database — without leaving the page. **Sample data** and **Cheat sheet** give example rows and a quick SQL reference.

The screenshot displays the XData-Web dashboard for a student named Max Schüler. The dashboard includes a header with the user's name and a sidebar with navigation options. The main content area is divided into several sections:

- YOUR LEARNING DASHBOARD:** Shows the student's name, level (Level 6), and progress (256 / 3455 XP). It also displays a 5-day streak and a score of 100%.
- LEADERBOARD:** Shows the student's rank (42 / 2) and top 10% score. The student is currently at the top of the leaderboard with 1 solved problem and 10 pts.
- RADDES:** A section with various badges and achievements, including 'First Solution', 'Diligent', 'SQL Pro', 'Perfectionist', 'Climber', and 'On a Roll'.
- MY ASSIGNMENTS:** Shows the current assignment 'Woche 1: SELECT Basics' with a progress bar and a score of 100%.
- QUESTIONS - WOCHE 1: SELECT BASICS:** A section for the current assignment, showing a problem statement and a hint.
- SCHEMA - UNIVERSITY DB:** A detailed view of the database schema, including tables, columns, and foreign-key relationships.
- LAST RESULT:** A section showing the student's overall score (3.0 / 3.0) and a comparison of their solution with the model solution.

The **SCHEMA - UNIVERSITY DB** section is expanded, showing the following tables and columns:

TABLE	COLUMNS	TABLE	COLUMNS	TABLE	COLUMNS
STUDENTS	id (PRIMARY KEY), name, age, major, gpa	COURSES	course_id (PRIMARY KEY), title, credits	ENROLLMENTS	st_id (FOREIGN KEY), course_id (FOREIGN KEY), grade

The **LAST RESULT** section shows the student's overall score (3.0 / 3.0) and a comparison of their solution with the model solution. The comparison table is as follows:

COMPONENT	YOUR SOLUTION	MODEL SOLUTION
Tables (FROM)	✓ STUDENTS (1+1)	STUDENTS
Columns (SELECT)	✓ name (1+1)	name
Conditions (WHERE)	✓ age > 22 (1+1)	age > 22
Joins (JOIN)	-	-
Grouping (GROUP BY)	-	-
Sorting (ORDER BY)	-	-
Filter (HAVING)	-	-
Duplicates (DISTINCT)	✓ No (1+1)	NO

Figure 20: The *Schema* helper expands the database structure (tables, columns, keys and foreign-key relationships) inline.

Hints. If you are stuck, reveal hints step by step with **Show hint**. Hints are optional and are unlocked one at a time.




Figure 21: Hints are revealed one step at a time.

5.4 Understanding the result

After submitting, the **Last result** panel shows an overall score and a **component-by-component** comparison of your query against the model solution — tables, columns, conditions, joins, grouping, sorting, filters and duplicates. Each matching part earns partial credit, so a near-miss still scores. You can **Compare expected output** or, if you believe the grade is wrong, **Request a regrade** (the instructor reviews it under *Regrade requests*).

6 SQL Lab (Diagnostic Lab)

The **SQL Lab** is a grading *simulator*. Pick a schema, enter a model solution and a candidate query, optionally adjust the grading weights, and run a structural evaluation — the same engine that grades student submissions. It is ideal for instructors testing a question or for anyone exploring how the partial marking works.

The screenshot displays the SQL Lab (Diagnostic Lab) interface. At the top, there is a navigation bar with links to 'XData /sql', 'Dashboard', 'SQL Lab', and 'Dataset Playground'. The 'SQL Lab' link is active. On the right side of the navigation bar, there is a user profile for 'Admin' and a settings icon. Below the navigation bar, the main header area shows 'SQL Diagnostic Lab' with a subtitle 'Structural grading simulation for SQL solutions.' and a dropdown menu for 'students_1781727484'. A button labeled 'Adjust grading weights' is located below the header. The main content area is divided into two panels. The left panel, titled 'Model solution' with an 'INSTRUCTOR VIEW' tab, contains a code editor with the following SQL query:

```
1 SELECT user_name FROM xdata_users WHERE
   internal_user_id = 'admin-uuid-admin1';
```

 The right panel, titled 'Student query' with a 'TEST CANDIDATE' tab, contains a code editor with the following SQL query:

```
1 SELECT user_name FROM xdata_users;
```

 Below these panels is a large blue button labeled 'Run structural evaluation'. At the bottom of the interface, there is a footer with the text 'XData IIT Bombay · Automated SQL Grading' on the left and '© 2026 — built for learning, query by query.' on the right.

Figure 22: SQL Lab: a model solution (left) and a student/candidate query (right).

The screenshot displays the SQL Diagnostic Lab interface. At the top, there's a navigation bar with 'XData /sql', 'Dashboard', 'SQL Lab', and 'Dataset Playground'. The user is logged in as 'Admin'. The main heading is 'SQL Diagnostic Lab' with a subtitle 'Structural grading simulation for SQL solutions.' and a student ID 'students_1781727484'. Below this is a section to 'Adjust grading weights'. The interface is split into two panels: 'Model solution' (Instructor View) and 'Student query' (Test Candidate). The 'Model solution' shows a query: `SELECT user_name FROM xdata_users WHERE internal_user_id = 'admin-uuid-admin1';`. The 'Student query' shows a query: `SELECT user_name FROM xdata_users;`. A large blue button labeled 'Run structural evaluation' is positioned below the queries. The 'STRUCTURAL ANALYSIS & FEEDBACK' section shows an 'OVERALL SCORE' of 2.0 / 3.0 (67%). Below this is a table comparing the 'MAIN QUERY' components between 'YOUR SOLUTION' and 'MODEL SOLUTION'.

COMPONENT	YOUR SOLUTION	MODEL SOLUTION
Tables (FROM)	✓ XDATA_USERS (+1.0)	XDATA_USERS
Columns (SELECT)	✓ user_name (+1.0)	user_name
Conditions (WHERE)	-	internal_user_id = 'admin-uuid-admin1'
Joins (JOIN)	-	-
Grouping (GROUP BY)	-	-
Sorting (ORDER BY)	-	-
Filter (HAVING)	-	-
Duplicates (DISTINCT)	✓ No +0.0	NO

At the bottom, there's a footer with 'XData IIT Bombay · Automated SQL Grading' and '© 2026 — built for learning, query by query.'

Figure 23: SQL Lab result: an overall score plus a per-component structural breakdown.

7 Dataset Playground

The **Dataset Playground** generates targeted test datasets that expose errors (“mutants”) in a SQL query, using the SMT solver (Z3). Select a schema, paste a reference query, choose which **mutation types** to target, and click **Generate dataset**. The generated INSERT statements appear on the right and can be copied or downloaded as .sql.

Dataset Playground

Generate targeted datasets to detect mutants of your SQL query.

SELECT DATABASE SCHEMA

students_1781727484

REFERENCE QUERY (SQL)

```
1 e.g. SELECT name FROM students WHERE age > 20;
```

MUTATION TYPES

☒ Selection ☒ Joins (Equi) ☐ Joins (Non-Equi)

☒ Aggregation ☐ Distinct ☐ Group By

☐ Having

Generate dataset

GENERATED DATASET (SQL INSERTS)

No dataset generated yet.
Click "Generate dataset".

XData IIT Bombay · Automated SQL Grading

© 2026 — built for learning, query by query.

Figure 24: Dataset Playground: choose a schema and mutation types, then generate a discriminating dataset.

Important — qualify your columns. The reference query must use *table-qualified* column names, e.g. `SELECT employees.name FROM employees WHERE employees.salary > 50000` (not just `name/salary`). This is a requirement of the underlying data-generation engine.

A successful run produces a dataset that includes *boundary* values — exactly the rows needed to tell a correct query apart from a subtly wrong one. For `employees.salary > 50000`, the engine generates rows with `salary = 50001` (passes), `50000` (the boundary that kills a `>=` mutant) and a non-matching value:

The screenshot displays the XData Dataset Playground interface. At the top, there's a navigation bar with links to XData /sql, Dashboard, SQL Lab, and Dataset Playground. The main header shows the 'Dataset Playground' title and a 'Load schema (.sql)' button. Below this, the interface is divided into two main sections.

Left Panel:

- SELECT DATABASE SCHEMA:** A dropdown menu showing 'Company DB'.
- REFERENCE QUERY (SQL):** A text area containing the query: `1 SELECT employees.name FROM employees WHERE employees.salary >`.
- MUTATION TYPES:** A section with checkboxes for various mutation types:
 - ☒ Selection
 - ☒ Joins (Equi)
 - ☐ Joins (Non-Equi)
 - ☒ Aggregation
 - ☐ Distinct
 - ☐ Group By
 - ☐ Having
- Generate dataset:** A blue button with a play icon.

Right Panel:

- GENERATED DATASET (SQL INSERTS):** A text area showing the generated SQL inserts:


```
insert into DEPARTMENTS values ('3',null,'4')
insert into DEPARTMENTS values ('5',null,'6')
insert into EMPLOYEES values ('7',null,'8','50001',null)
insert into DEPARTMENTS values ('1',null,'21240')
insert into DEPARTMENTS values ('-92279',null,'3')
insert into EMPLOYEES values ('4',null,'5','-97560',null)
insert into DEPARTMENTS values ('3',null,'4')
insert into DEPARTMENTS values ('5',null,'6')
insert into EMPLOYEES values ('7',null,'8','50000',null)
```
- 9 test data row(s) generated.** A green status message at the bottom of the right panel.

At the bottom of the interface, there's a footer with 'XData IIT Bombay · Automated SQL Grading' on the left and '© 2026 — built for learning, query by query.' on the right.

Figure 25: A successful generation on a custom schema. Z3 produced boundary values (50000 / 50001) that discriminate the query from its mutants.

7.1 Mutation types

Each ticked **mutation type** adds one targeted “killing” dataset aimed at a specific class of mistake — *Selection* (WHERE), *Joins (equi / non-equi)*, *Aggregation*, *Distinct*, *Group By* and *Having*. Tick only the types relevant to your query: a type that does not apply (e.g. *Aggregation* for a query without a GROUP BY) simply adds no rows. Toggling a type changes the output — for the salary query above, enabling *Selection* adds the boundary-killing rows (9 rows total), while disabling it falls back to the base dataset (3 rows).




Figure 26: Mutation types are applied live: unticking a type removes its targeted rows from the generated dataset.

8 Tips & troubleshooting

- **Light / dark mode** — toggle with the sun/moon icon in the top bar (Figure 27).
- **Change password** — the key icon in the top bar opens the change-password dialog (Figure 28).
- **“No assignment / course selected”** — some panels (Statistics, Schema management) need you to pick a course or assignment first; this is expected.
- **Ports** — the app uses 80 (frontend), 8080 (backend) and 5433 (database). Make sure they are free.
- **Fresh start** — `docker compose down -v` wipes the database and re-seeds the `admin1` account on the next start.

The screenshot shows the XData Admin Center interface in light mode. The top navigation bar includes the XData logo, a search bar, and links to Dashboard, SQL Lab, and Dataset Playground. The user is logged in as Admin. The main section is titled 'Admin Center' with the subtitle 'System-wide management & monitoring.' Below this, there are tabs for 'Users' and 'Courses'. The 'Users' tab is active, showing 'User management' options: 'Template' and 'Import CSV'. A 'Create new user' form is displayed with fields for 'FULL NAME' (e.g. John Doe), 'LOGIN ID' (e.g. jdoe), 'EMAIL ADDRESS' (john@example.com), 'PASSWORD' (masked), and 'ROLE' (STUDENT). There is also a 'COURSE ASSIGNMENT' section with a checkbox for 'Datenbanken I'. A 'Create user' button is at the bottom right of the form. Below the form, there is a section for 'Assigned users' with a search bar and a 'Newest first' sort option. A table of assigned users is visible below this section.

Figure 27: The interface in light mode (toggled from the top bar).

The screenshot shows the XData Admin Center interface in dark mode. A 'Change password' dialog is open in the center of the screen. The dialog has a title bar with a key icon and a close button. It contains the text 'Choose a new password (at least 8 characters).' and three input fields: 'CURRENT PASSWORD', 'NEW PASSWORD', and 'CONFIRM NEW PASSWORD'. At the bottom of the dialog are 'Cancel' and 'Change password' buttons. The background shows the 'Assigned users' table from the previous figure, but it is dimmed.

Figure 28: The change-password dialog (key icon, top bar).

XData IIT Bombay · Automated SQL Grading — built for learning, query by query.